

Umang Gupta

✉ umanggup@usc.edu — 🏠 umgupta.github.io — 🌐 github.com/umgupta

Education

University of Southern California (USC), Los Angeles Ph.D. in Computer Science (<i>Expected May 2023</i>) Thesis: Controlling Information for Fairness and Privacy in Machine Learning Advisor: Prof. Greg Ver Steeg, Information Sciences Institute USC	2017 - 2023 GPA: 3.81/4
Indian Institute of Technology Delhi (IIT Delhi), India B. Tech & M. Tech (Dual Degree), Electrical Engineering (<i>Silver Medalist</i>) Thesis: Image Classification with Ontology and Deep Learning Advisor: Prof. Santanu Chaudhury, Indian Institute of Technology, Delhi	2010 - 2015 GPA: 8.9/10

Work Experience

Google Gemini, Mountain View <i>Research Engineer</i> <ul style="list-style-type: none">Developed data and training pipelines for Gemini suite of models, particularly focused on tool-use dataDeveloping novel techniques to improve tool-use capabilities of Gemini modelsReceived 5 spot and 6 peer bonuses within a year for delivering critical features in Gemini releases	Nov'23 - Present
Machine Learning Center of Excellence, Morgan Stanley, New York <i>Research Intern</i> <ul style="list-style-type: none">Contributed to several research projects focused on improving time-series forecastingAnalyzed quantization error in transformers, leading to an efficient vector-quantized transformer model whose complexity <i>scales linearly with sequence size</i> [Submitted to ICLR]	Jun'22 - Aug'22
Amazon Alexa-AI, Boston (Virtual) <i>Applied Scientist Intern</i> <ul style="list-style-type: none">Studied biases in text generated from language models, especially the effect of model size on biasesDevised knowledge-distillation approach for reducing biases in language models [ACL 2022]	May'21 - Aug'21
Futurewei Technnologies, Santa Clara <i>Research Intern</i> <p>Investigated reinforcement learning algorithms for end-to-end training of seq2seq dialogue models</p>	Jun'18 - Aug'18
Software Engineering Experiences <ul style="list-style-type: none">↔ Visa Inc., Bangalore, India <i>Senior Software Engineer (User-Interface and JavaScript Development)</i>↔ LCI Lab, University of British Columbia, Vancouver, Canada <i>MITACS Research Intern</i>↔ Amagi Media Labs, Bangalore, India <i>Software Development Intern</i>↔ Sohumi Innovation Labs, Delhi, India <i>Intern</i>	Aug'15 - Jul'17 May'14 - Jul'14 May'13 - Jul'13 May'12 - Dec'12

Selected Publications [\[Google Scholar\]](#)

Fair Machine Learning

- Umang Gupta**, Jwala Dhamala, ..., Rahul Gupta, Kai-Wei Chang, Greg Ver Steeg, Aram Galstyan
Equitable Text Generation with Distilled Language Models via Counterfactual Role Reversal
Findings of the Association for Computational Linguistics (ACL 2022)
- Umang Gupta**, Aaron Ferber, Bistra Dilkina, Greg Ver Steeg
Controllable Guarantees for Fair Outcomes via Contrastive Information Estimation
Proceedings of the 35th AAAI Conference on Artificial Intelligence (AAAI 2021)

Ninareh Mehrabi, **Umang Gupta**, Fred Morstatter, Greg Ver Steeg, Aram Galstyan
Attributing Fair Decisions with Attention Interventions
TrustNLP Workshop @ NAACL 2022

Private, and Distributed Learning

Umang Gupta, Dmitris Stripelis, Pradeep Lam, Paul Thompson, José Luis Ambite, Greg Ver Steeg
Membership Inference Attacks on Deep Regression Models for Neuroimaging
Proceedings of the 4th Conference on Medical Imaging with Deep Learning (MIDL 2021)

Dmitris Stripelis, **Umang Gupta**, Greg Ver Steeg, José Luis Ambite
Federated Progressive Sparsification (Purge, Merge, Tune)+
Federated Learning Workshop @ NeurIPS 2022

Others

Umang Gupta, Tamoghna Chattopadhyay, Nikhil Dhinagar, Paul Thompson, Greg Ver Steeg
Transferring Models Trained on Natural Images to 3D MRI via Position Encoded Slice Models
IEEE International Symposium on Biomedical Imaging (ISBI), 2023

Sahil Garg, **Umang Gupta**, Yu Chen, Syamantak Gupta, Yeshaya Adler, Anderson Schneider, Yuriy Nevmyvaka
Estimating Transfer Entropy under Long Ranged Dependencies
Uncertainty in Artificial Intelligence (UAI), 2022

Umang Gupta, Pradeep Lam, Greg Ver Steeg, Paul Thompson
Improved Brain Age Estimation with Slice-based Set Networks
IEEE International Symposium on Biomedical Imaging (ISBI), 2021

Nitin Kamra, **Umang Gupta**, Fei Fang, Yan Liu, Milind Tambe
Policy Learning for Continuous Space Security Games using Neural Networks
Proceedings of the 32nd AAAI Conference on Artificial Intelligence (AAAI 2018)

Nitin Kamra, **Umang Gupta**, Yan Liu
Deep Generative Dual Memory Network for Continual Learning
<https://arxiv.org/abs/1710.10368>

Selected Projects

Information-Theoretic Measures for Privacy & Fairness

Jul'19 - Jul'23

↪ *Controlling Fairness via Mutual Information Minimization* [AAAI 2021]

- Characterized an information-theoretic measure for statistical parity, a popular fairness measure
- Demonstrated an effective method for controlling fairness through contrastive mutual information estimation
- Showcased *better fairness-accuracy trade-offs* compared to recent competitive baselines

↪ *Minimizing Privacy Leakage during Training of the Neural Networks*

- Studied ways in which machine learning models may leak private training set information [MIDL 2021]
- Investigating and improving techniques for differential private training of neural networks [Ongoing]

Continual Learning with Neural Networks

Aug'17 - Jan'18

Best Theory Project Award, CSCI-599: Deep Learning, USC (2017)

- Taking inspiration from biological learning mechanisms in humans, we proposed dual generative (memory) models to solve the problem of catastrophic forgetting in Neural Networks
- Demonstrated *better retention of previously seen concepts* when training samples arrive in a non-iid fashion

Technical Skills

Programming

Languages Python, C, Shell scripting, JavaScript, Java

Libraries & Tools PyTorch (among the top 10 [answerers](#) on [StackOverflow](#)), JAX, TensorFlow, ReactJS, BackboneJS, NodeJS

Awards and Honors

Program	Recipient of IITD Semester Merit Award for 4 of 8 semesters, 2010 - 2014
Rank 1	Highest GPA in Dual Degree Program, Department of Electrical Engineering
Research	Best Theory Project Award for arxiv:1710.10368 , CSCI-599: Deep Learning, USC (2017)
Best Essay	Among the top 20 National Winners, International Year of Forest, 2011 certification program
Hackathons	First Prize, GS Quantify 2014, an annual computing competition organized by Goldman Sachs First Prize, Bing Hackathon 2016, a machine learning contest organized by Microsoft Bing
NIUS 2011	Among 30 students invited to research at HBCSE, Mumbai, under National Initiative for Undergraduate Science (NIUS) program
Competitive Exams	Recipient of KVPY scholarship 2010, Dept. of Science and Technology, Govt. of India Among the top 1% in Physics (NSEP); Astronomy (NSEA); Chemistry (NSEC) Olympiads 2009 Secured All India Rank 410 in IIT-Joint Entrance Exam among 0.5 million candidates

Professional Services

Reviewing Medical Imaging and Analysis (MedIA) Journal
International Conference on Learning Representations (ICLR) – 2023
Association for Computational Linguistics (ACL) – 2023
Uncertainty in Artificial Intelligence (UAI) – 2023
Asian Conference on Machine Learning (ACML) – 2022
Artificial Intelligence and Statistics (AISTATS) – 2021
Medical Imaging with Deep Learning – 2021, 2022, 2023
Workshops: {ENLSP, MedNeurIPS} @ NeurIPS 2022; TrustNLP @ NAACL 2022; FL4NLP @ ACL 2022;